

Some Helper Function

✓ Softmax function:

```
import numpy as np
def softmax(z):
    """
    Compute the softmax probabilities for a given input matrix.
    Parameters:
    z (numpy.ndarray): Logits (raw scores) of shape (m, n), where

    - m is the number of samples.
    - n is the number of classes.

    Returns:
    numpy.ndarray: Softmax probability matrix of shape (m, n), where
    each row sums to 1 and represents the probability
    distribution over classes.

    Notes:
    - The input to softmax is typically computed as:  $z = XW + b$ .
    - Uses numerical stabilization by subtracting the max value per row.
    """
    # Your Code Here.
    z = np.atleast_2d(z)

    z_max = np.max(z, axis=1, keepdims=True)
    exp_z = np.exp(z - z_max)
    softmax_probs = exp_z / np.sum(exp_z, axis=1, keepdims=True)
    return softmax_probs
```

✓ Softmax Test case

This test case checks that each row in the resulting softmax probabilities sums to 1, which is the fundamental property of softmax

```
def test_softmax():
    # Test case setup
    z_test = np.array([[2.0, 1.0, 0.1], [1.0, 1.0, 1.0]])

    # Calculate softmax
    softmax_output = softmax(z_test)

    # Calculate rows sum
    row_sums = np.sum(softmax_output, axis=1)
    # Verify results
    try:
        np.allclose(row_sums, 1)
        print(f"Test failed: Row sum are {row_sums}")
        print("Softmax Function passed the test case")
    except AssertionError as e:
        print((f"X Test failed: Row sums are {row_sums}"))
```

```
test_softmax()
```

```
↗ Test failed: Row sum are [1. 1.]
Softmax Function passed the test case
```

```
def predict_softmax(X, W, b):
    """
    Predict the class labels for a set of samples using the trained softmax model.
    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the
    number of features.
    W (numpy.ndarray): Weight matrix of shape (d, c), where c is the number of classes.
    b (numpy.ndarray): Bias vector of shape (c,).
    Returns:
```

```

numpy.ndarray: Predicted class labels of shape (n,), where each value is the index of the
predicted class.
"""
logist = np.dot(X,W)+b #Compute raw scores (logits)
probabilities = softmax(logist)
predicted_classes = np.argmax(probabilities,axis=1)
return predicted_classes

```

✓ Test Function for Prediction Function:

The test function ensures that the predicted class labels have the same number of elements as the input samples, verifying that the model produces a valid output shape.

```

# The test function ensures that the predicted class labels have the same number of elements as the
# input samples, verifying that the model produces a valid output shape.
# Define test case
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)
# Expected Output:
# The function should return an array with class labels (0, 1, or 2)
y_pred_test = predict_softmax(X_test, W_test, b_test)
# Validate output shape
assert y_pred_test.shape == (3,), f"Test failed: Expected shape (3,), got {y_pred_test.shape}"
# Print the predicted labels
print("Predicted class labels:", y_pred_test)

```

↻ Predicted class labels: [1 1 0]

✓ 1. Implementation of Loss Function:

```

def loss_softmax(y_pred, y):
    """
    Compute the cross-entropy loss for a single sample.
    Parameters:
    y_pred (numpy.ndarray): Predicted probabilities of shape (c,) for a single sample,
    where c is the number of classes.

    y (numpy.ndarray): True labels (one-hot encoded) of shape (c,), where c is the number of classes.
    Returns:
    float: Cross-entropy loss for the given sample.
    """
    epsilon = 1e-12 #small constant to prevent log(0)
    y_pred = np.clip(y_pred,epsilon,1.0 - epsilon) # Clip values to prevent numerical instability
    loss = -np.sum(y*np.log(y_pred))
    return loss

```

✓ Test case for Loss Function:

This test case Compares loss for correct vs. incorrect predictions.

- Expects low loss for correct predictions.
- Expects high loss for incorrect predictions.

```

import numpy as np
# This test case Compares loss for correct vs. incorrect predictions.
# Expects low loss for correct predictions.
# Expects high loss for incorrect predictions.
# Define correct predictions (low loss scenario)
y_true_correct = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True one-hot labels
y_pred_correct = np.array([[0.9, 0.05, 0.05],
[0.1, 0.85, 0.05],
[0.05, 0.1, 0.85]]) # High confidence in the correct class

# Define incorrect predictions (high loss scenario)
y_pred_incorrect = np.array([[0.05, 0.05, 0.9], # Highly confident in the wrong class
[0.1, 0.05, 0.85],

```

```
[0.85, 0.1, 0.05]])
```

```
# Compute loss for both cases
loss_correct = loss_softmax(y_pred_correct, y_true_correct)
loss_incorrect = loss_softmax(y_pred_incorrect, y_true_correct)
# Validate that incorrect predictions lead to a higher loss
assert loss_correct < loss_incorrect, f"Test failed: Expected loss_correct < loss_incorrect, but got{loss_correct:.4f} >= {loss_incorrect:.4f}"
# Print results
print(f"Cross-Entropy Loss (Correct Predictions): {loss_correct:.4f}")
print(f"Cross-Entropy Loss (Incorrect Predictions): {loss_incorrect:.4f}")
```

```
↗ Cross-Entropy Loss (Correct Predictions): 0.4304
Cross-Entropy Loss (Incorrect Predictions): 8.9872
```

✓ 2. Implementation of Cost Function:

```
def cost_softmax(X, y, W, b):
    """
    Compute the average softmax regression cost (cross-entropy loss) over all samples.
    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the
    number of features.
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c), where n is the number of
    samples and c is the number of classes.
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).
    Returns:
    float: Average softmax cost (cross-entropy loss) over all samples.
    """
    # Compute logits (Linear Transformation)
    logits = np.dot(X,W)+b
    # Compute predicted probabilities using softmax
    predictions = softmax(logits)
    # Calculate cross-entropy loss
    cost = loss_softmax(predictions,y)
    return cost
```

✓ Test Case for Cost Function:

The test case assures that the cost for the incorrect prediction should be higher than for the correct prediction, confirming that the cost function behaves as expected.

```
import numpy as np

# Example 1: Correct Prediction (Closer predictions)
X_correct = np.array([[1.0, 0.0], [0.0, 1.0]]) # Feature matrix for correct predictions
y_correct = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, matching predictions)
W_correct = np.array([[5.0, -2.0], [-3.0, 5.0]]) # Weights for correct prediction
b_correct = np.array([0.1, 0.1]) # Bias for correct prediction

# Example 2: Incorrect Prediction (Far off predictions)
X_incorrect = np.array([[0.1, 0.9], [0.8, 0.2]]) # Feature matrix for incorrect predictions
y_incorrect = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, incorrect predictions)
W_incorrect = np.array([[0.1, 2.0], [1.5, 0.3]]) # Weights for incorrect prediction
b_incorrect = np.array([0.5, 0.6]) # Bias for incorrect prediction


# Compute cost for correct predictions
cost_correct = cost_softmax(X_correct, y_correct, W_correct, b_correct)

# Compute cost for incorrect predictions
cost_incorrect = cost_softmax(X_incorrect, y_incorrect, W_incorrect, b_incorrect)

# Check if the cost for incorrect predictions is greater than for correct predictions
assert cost_incorrect > cost_correct, f"Test failed: Incorrect cost {cost_incorrect} is not greater than correct cost {cost_correct}"

# Print the costs for verification
print("Cost for correct prediction:", cost_correct)
print("Cost for incorrect prediction:", cost_incorrect)

print("Test passed!")
```

 Cost for correct prediction: 0.0012468728266698647
Cost for incorrect prediction: 0.5986172271889223
Test passed!

✓ Implementation of Compute Gradients

```
def compute_gradient_softmax(X, y, W, b):
    """
    Compute the gradients of the cost function with respect to weights and biases.
    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).
    Returns:
    tuple: Gradients with respect to weights (d, c) and biases (c,).
    """
    n,d = X.shape
    z= np.dot(X,W)+b
    y_pred = softmax(z)
    grad_W = np.dot(X.T, (y_pred-y))/n # Calculate with respect to weights
    grad_b = np.sum(y_pred-y, axis=0) /n #Calculate with respect to biases
    return grad_W, grad_b
```

✓ Test case for compute_gradient function:

The test checks if the gradients from the function are close enough to the manually computed gradients using `np.allclose`, which accounts for potential floating-point discrepancies.

```
import numpy as np

# Define a simple feature matrix and true labels
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)
y_test = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True labels (one-hot encoded, 3 classes)

# Define weight matrix and bias vector
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

# Compute the gradients using the function
grad_W, grad_b = compute_gradient_softmax(X_test, y_test, W_test, b_test)
```

```
# Manually compute the predicted probabilities (using softmax function)
z_test = np.dot(X_test, W_test) + b_test
y_pred_test = softmax(z_test)

# Compute the manually computed gradients
grad_W_manual = np.dot(X_test.T, (y_pred_test - y_test)) / X_test.shape[0]
grad_b_manual = np.sum(y_pred_test - y_test, axis=0) / X_test.shape[0]

# Assert that the gradients computed by the function match the manually computed gradients
assert np.allclose(grad_W, grad_W_manual), f"Test failed: Gradients w.r.t. W are not equal.\nExpected: {grad_W_manual}\nGot: {grad_W}"
assert np.allclose(grad_b, grad_b_manual), f"Test failed: Gradients w.r.t. b are not equal.\nExpected: {grad_b_manual}\nGot: {grad_b}"

# Print the gradients for verification
print("Gradient w.r.t. W:", grad_W)
print("Gradient w.r.t. b:", grad_b)

print("Test passed!")
```

```
→ Gradient w.r.t. W: [[ 0.1031051  0.01805685 -0.12116196]
 [-0.13600547  0.00679023  0.12921524]]
Gradient w.r.t. b: [-0.03290036  0.02484708  0.00805328]
Test passed!
```

✓ Implementing Gradient Descent:

```
def gradient_descent_softmax(X,W,y,b,alpha,n_iter,show_cost=False):
    """
    Perform gradient descent to optimize the weights and biases.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).
    alpha (float): Learning rate.
    n_iter (int): Number of iterations.
    show_cost (bool): Whether to display the cost at intervals.

    Returns:
    tuple: Optimized weights, biases, and cost history.
    """

    cost_history = []

    for i in range(n_iter):
        grad_W,grad_b = compute_gradient_softmax(X,y,W,b)

        # Updating Weights and biases using the gradients
        W = W- alpha * grad_W
        b = b- alpha * grad_b

        # Compute and store cost
        cost = cost_softmax(X,y,W,b)
        cost_history.append(cost)

        if show_cost and (i%100==0 or i==n_iter-1):
            print(f"Iteration {i}: Cost = {cost:.6f}")

    return W,b, cost_history
```

✓ Preparing Dataset:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def load_and_prepare_mnist(csv_file,test_size=0.2,random_state=42):
    """
    Reads the MNIST CSV file, splits data into train/test sets, and plots one image per class.
```

```

Arguments:
csv_file (str)      : Path to the CSV file containing MNIST data.
test_size (float)   : Proportion of the data to use as the test set (default: 0.2).
random_state (int)  : Random seed for reproducibility (default: 42).

Returns:
X_train, X_test, y_train, y_test : Split dataset.
"""

# Load dataset
df = pd.read_csv(csv_file)

# Separate labels and features
y = df.iloc[:,0].values #First column is the label
x = df.iloc[:,1:].values # Remaining columns are pixel values

# Noramlize pixels values (optional but recommend)
X = x/255.0

# Split data into train and test sets

X_train, X_test, y_train,y_test = train_test_split(X,y,test_size=test_size,random_state=random_state)

# Plot one sample image per class
plot_sample_images(X,y)

return X_train,X_test,y_train,y_test

def plot_sample_images(X,y):
    """
    Plots one sample image for each digit class (0-9).

    Arguments:
    X (np.ndarray): Feature matrix containing pixel values.
    y (np.ndarray): Labels corresponding to images.
    """
    plt.figure(figsize=(10,4))
    unique_classes = np.unique(y) # Get unique class labels

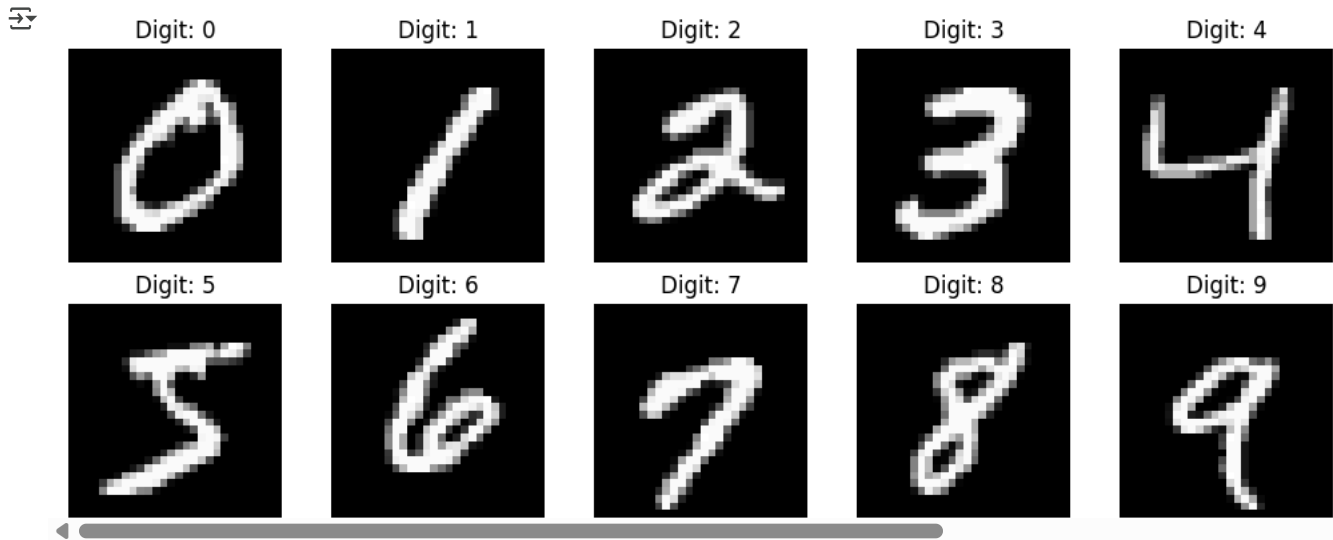
    for i, digit in enumerate(unique_classes):
        index = np.where(y==digit)[0][0] # Find first occurance of the class
        image = X[index].reshape(28,28) # Reshape 1D array to 28x28

        plt.subplot(2,5,i+1)
        plt.imshow(image,cmap='gray')
        plt.title(f'Digit: {digit}')
        plt.axis('off')

    plt.tight_layout()
    plt.show()

csv_file = '/content/drive/MyDrive/Artificial_Intelligence/Week2/Worksheet2/mnist_dataset.csv'
X_train,X_test,y_train,y_test = load_and_prepare_minst(csv_file)

```



✓ A Quick debugging Step:

```
# Assert that X and y have matching lengths
assert len(X_train) == len(y_train), f"Error: X and y have different lengths! X={len(X_train)} , y={len(y_train)}"
print("Move forward: Dimension of Feture Matrix X and label vector y matched.")
```

➡ Move forward: Dimension of Feture Matrix X and label vector y matched.

✓ Train the Model:

```
print(f"Training data shape: {X_train.shape}")
print(f"Test data shape: {X_test.shape}")
```

➡ Training data shape: (48000, 784)
Test data shape: (12000, 784)

```
from sklearn.preprocessing import OneHotEncoder
```

```
# Check if y_train is one-hot encoded
if len(y_train.shape) == 1:
    encoder = OneHotEncoder(sparse_output=False) # Use sparse_output=False for newer versions of sklearn
    y_train = encoder.fit_transform(y_train.reshape(-1, 1)) # One-hot encode labels
    y_test = encoder.transform(y_test.reshape(-1, 1)) # One-hot encode test labels
```

```
# Now y_train is one-hot encoded, and we can proceed to use it
d = X_train.shape[1] # Number of features (columns in X_train)
c = y_train.shape[1] # Number of classes (columns in y_train after one-hot encoding)
```

```
# Initialize weights with small random values and biases with zeros
W = np.random.randn(d, c) * 0.01 # Small random weights initialized
b = np.zeros(c) # Bias initialized to 0
```

```
# Set hyperparameters for gradient descent
alpha = 0.1 # Learning rate
n_iter = 1000 # Number of iterations to run gradient descent
```

```
# print(f'X shape: {X_train.shape}, W shape: {W.shape}, b shape: {b.shape}')
```


```
# Train the model using gradient descent
W_opt, b_opt, cost_history = gradient_descent_softmax(X_train, W, y_train, b, alpha, n_iter, show_cost=True)
```

```
# Plot the cost history to visualize the convergence
plt.plot(cost_history)
plt.title('Cost Function vs. Iterations')
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.grid(True)
plt.show()
```

```

↻ Iteration 0: Cost = 105034.594116
Iteration 100: Cost = 29162.110983
Iteration 200: Cost = 23508.684417
Iteration 300: Cost = 21173.055160
Iteration 400: Cost = 19823.680666
Iteration 500: Cost = 18917.027185
Iteration 600: Cost = 18252.888972
Iteration 700: Cost = 17738.443725
Iteration 800: Cost = 17324.089057
Iteration 900: Cost = 16980.610905
Iteration 999: Cost = 16692.231772

```



✓ Evaluating the model

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score

# Evaluation Function
def evaluate_classification(y_true,y_pred):
    """
    Evaluate classification performance using confusion matrix, precision, recall, and F1-score.

    Parameters:
    y_true (numpy.ndarray): True labels
    y_pred (numpy.ndarray): Predicted labels

    Returns:
    tuple: Confusion matrix, precision, recall, F1 score
    """

    # Compute confusion matrix
    cm = confusion_matrix(y_true,y_pred)

    # Compute precision, recall , and F1-score
    precision = precision_score(y_true,y_pred,average='weighted')
    recall = recall_score(y_true,y_pred,average='weighted')
    f1 = f1_score(y_true,y_pred,average='weighted')

    return cm,precision,recall,f1

# Prdict on the test score
y_pred_test = predict_softmax(X_test,W_opt,b_opt)

# Evaluate accuracy
y_test_labels = np.argmax(y_test,axis=1) # True labels in numeric form

# Evaluate the model
cm,precision,recall,f1 = evaluate_classification(y_test_labels,y_pred_test)

```



```

# Print the evaluation metrics
print("\nConfusion Matrix:")
print(cm)
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-score: {f1:.2f}")

# Visualizing the Confusion Matrix
fig, ax = plt.subplots(figsize=(12,12))
cax = ax.imshow(cm,cmap="Blues") # Use a color map for better visualization

# Dynamic number of classes
num_classes = cm.shape[0]
ax.set_xticks(range(num_classes))
ax.set_yticks(range(num_classes))
ax.set_xticklabels([f'Predicted {i}' for i in range(num_classes)])
ax.set_yticklabels([f'Actual {i}' for i in range(num_classes)])

# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j,i,cm[i,j], ha='center', va='center', color='white' if cm[i,j] > np.max(cm) / 2 else 'black')

# Add labels to each cell in the confusion matrix
ax.grid(False)
plt.title('Confusion Matrix', fontsize=14)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('Actual Label', fontsize=12)

# Adjust layout
plt.tight_layout()
plt.colorbar(cax)
plt.show()

```



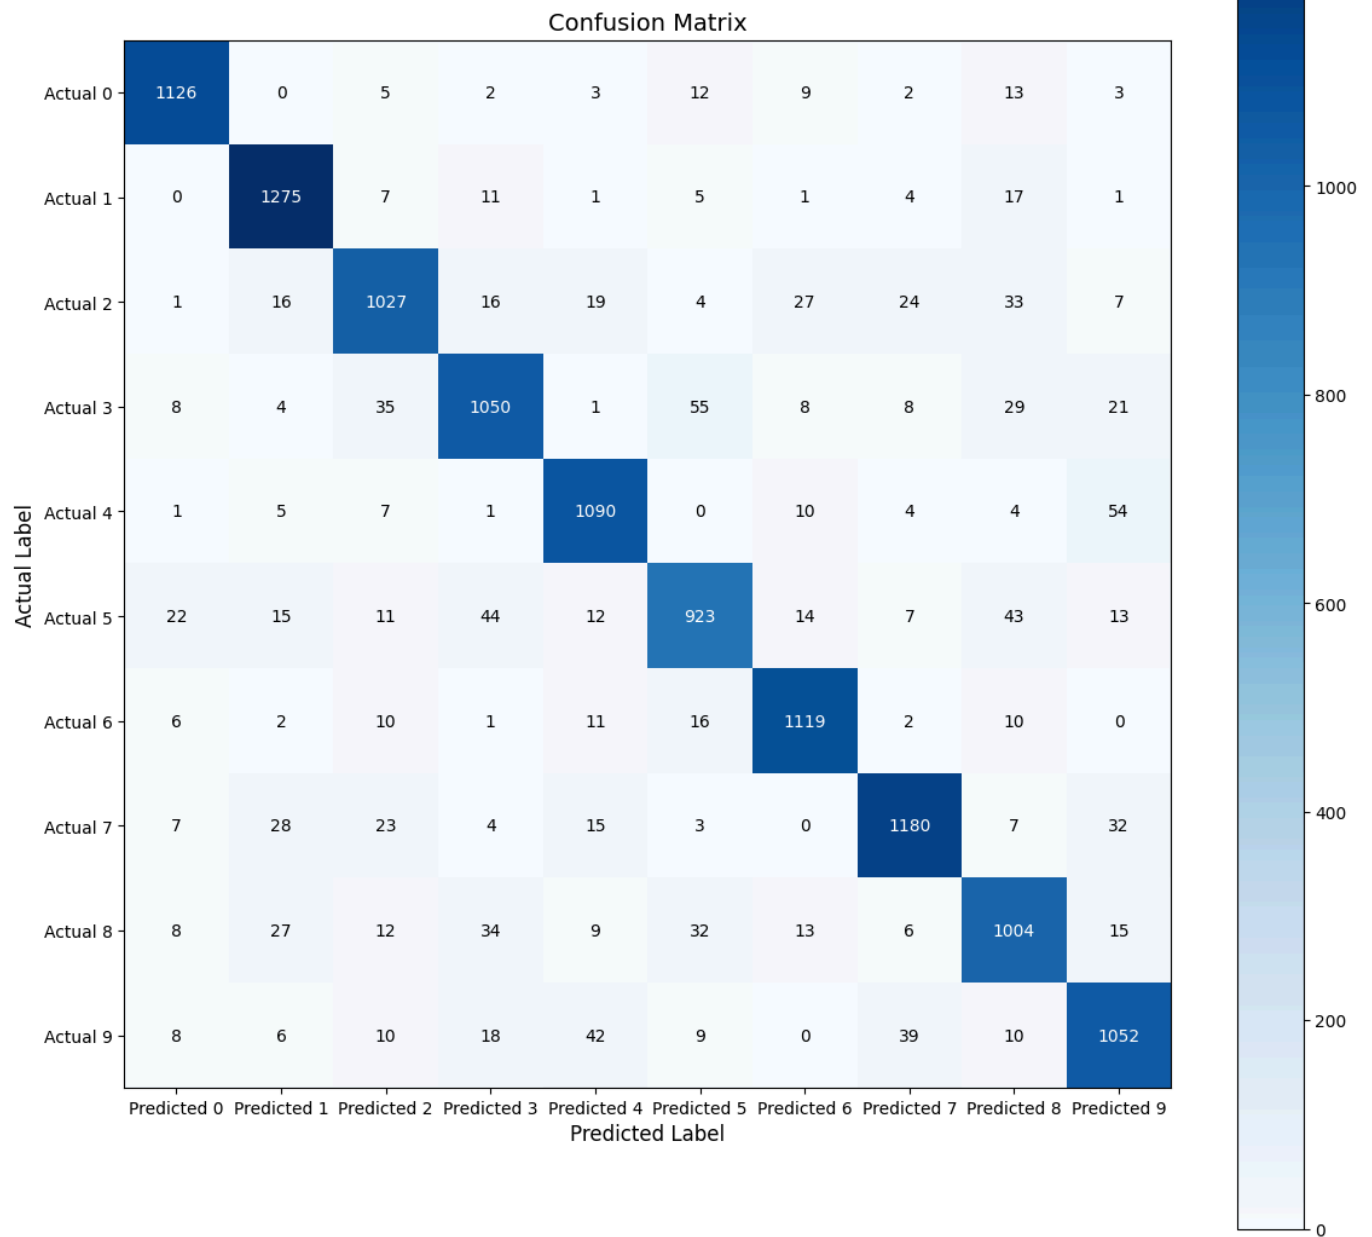
Confusion Matrix:

```
[[1126  0  5  2  3 12  9  2 13  3]
 [  0 1275  7 11  1  5  1  4 17  1]
 [  1 16 1027 16 19  4 27 24 33  7]
 [  8  4 35 1050  1 55  8  8 29 21]
 [  1  5  7  1 1090  0 10  4  4 54]
 [ 22 15 11 44 12 923 14  7 43 13]
 [  6  2 10  1 11 16 1119  2 10  0]
 [  7 28 23  4 15  3  0 1180  7 32]
 [  8 27 12 34  9 32 13  6 1004 15]
 [  8  6 10 18 42  9  0 39 10 1052]]
```

Precision: 0.90

Recall: 0.90

F1-score: 0.90



```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification, make_circles
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

# Set random seed for reproducibility
np.random.seed(42)

# Generate linearly separable dataset
X_linear_separable, y_linear_separable = make_classification(n_samples=200, n_features=2,
n_informative=2,
n_redundant=0, n_clusters_per_class=1,
random_state=42)

# Split the data into training and testing sets
X_train_linear, X_test_linear, y_train_linear, y_test_linear = train_test_split(
X_linear_separable, y_linear_separable, test_size=0.2, random_state=42
)

# Train logistic regression model on linearly separable data
logistic_model_linear_separable = LogisticRegression()
logistic_model_linear_separable.fit(X_train_linear, y_train_linear)

# Generate non-linearly separable dataset (circles)
X_non_linear_separable, y_non_linear_separable = make_circles(n_samples=200, noise=0.1, factor=0.5,
random_state=42)

# Split the data into training and testing sets
X_train_non_linear, X_test_non_linear, y_train_non_linear, y_test_non_linear = train_test_split(
X_non_linear_separable, y_non_linear_separable, test_size=0.2, random_state=42
)

# Train logistic regression model on non-linearly separable data
logistic_model_non_linear_separable = LogisticRegression()
logistic_model_non_linear_separable.fit(X_train_non_linear, y_train_non_linear)

# Plot decision boundaries for linearly and non-linearly separable data
def plot_decision_boundary(ax, model, X, y, title):
    h = .02 # step size in the mesh
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    ax.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    ax.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Paired) # Corrected here
    ax.set_title(title)
    ax.set_xlabel('Feature 1') # Corrected here
    ax.set_ylabel('Feature 2') # Corrected here

# Create subplots
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# Plot decision boundary for linearly separable data (Training)
plot_decision_boundary(axes[0, 0], logistic_model_linear_separable, X_train_linear, y_train_linear,
'Linearly Separable Data (Training)') # Corrected here

# Plot decision boundary for linearly separable data (Testing)
plot_decision_boundary(axes[0, 1], logistic_model_linear_separable, X_test_linear, y_test_linear,
'Linearly Separable Data (Testing)') # Corrected here

# Plot decision boundary for non-linearly separable data (Training)
plot_decision_boundary(axes[1, 0], logistic_model_non_linear_separable, X_train_non_linear, y_train_non_linear,
'Non-linearly Separable Data (Training)') # Corrected here

```